



African Language Safety Fragility & Deliberative Reinforcement Learning Suite


(a) Cross-Lingual Refusal Fragility on Unsafe Prompts



(b) Layer-Wise Refusal Probability Drop RPD_ℓ



(c) Inference-Time VM-MCTS Search Safety Compliance



(d) Jacobian Lens Subspace Awakening Gain ($\|\delta\|_2 \leq 5.0$)

